

Zylo AI Readiness Sprint

McKinsey-Standard Upgrade Pack

10 Implemented Systems, Templates & Frameworks

Zylo Solutions | Internal Use Only | April 2026

How to use this document: This document implements every element identified in the McKinsey-standard gap analysis. Each section below is a working, ready-to-use system, template, or framework — not a recommendation. Read, customise the [bracketed fields], and deploy before Sprint 1.

ELEMENT 1 — The Zylo AI Readiness Matrix (The ZARM)

What this is: This is Zylo's proprietary scoring framework. It replaces the unnamed 4-dimensional assessment with a branded, repeatable methodology. Reference it by name in every client session, every deliverable, and all marketing materials.

Framework Definition

The Zylo AI Readiness Matrix (The ZARM) is a 4-axis scoring model used to assess a company's readiness to adopt, implement, and scale AI across its operations. Developed through Zylo's engagements across SMBs and growth-stage companies in North America, Europe, and the GCC, the Matrix evaluates every business process against four dimensions that together determine both the feasibility and the value of AI implementation.

The Four Axes — Definitions & Scoring Guide

Axis 1: Process Maturity

Scoring Question: How documented, consistent, and measurable is this process today?

- 1 — Chaotic:** Process is informal, undocumented, inconsistent. Different people do it differently every time.
- 2 — Developing:** Process is partially documented. Some consistency but relies on individual knowledge.
- 3 — Defined:** Process is documented. Most team members follow it. Some exceptions exist.
- 4 — Managed:** Process is fully documented, consistently followed, and has measurable outputs.
- 5 — Optimised:** Process is fully standardised, measured, and continuously improved.

Axis 2: Data Readiness

Scoring Question: Does the data required to train or run an AI model for this process exist, and is it accessible?

- 1 — No Data:** No digital data exists for this process. It is entirely manual with no record trail.
- 2 — Sparse:** Some data exists but is incomplete, inconsistent, or in non-digital formats.
- 3 — Partial:** Relevant data exists digitally but is siloed, unstructured, or requires significant cleaning.
- 4 — Available:** Clean, structured data exists and is accessible via existing systems or exports.
- 5 — Rich:** Extensive, clean, well-labelled historical data exists. Ready to use immediately.

Axis 3: Business Impact

Scoring Question: If this process were automated or enhanced by AI, what would the measurable business impact be?

- 1 — Negligible:** Automating this process would save minimal time or cost. Low strategic relevance.

- 2 — **Low:** Modest time or cost saving. Useful but not transformative.
- 3 — **Moderate:** Meaningful saving — e.g. 3-8 hours per week or a clearly measurable quality improvement.
- 4 — **High:** Significant saving — e.g. 10+ hours per week, major cost reduction, or direct revenue impact.
- 5 — **Critical:** Transformative. Automating this changes how the business operates and competes.

Axis 4: Integration Complexity

Scoring Question: How complex would it be to integrate an AI solution into this process given the current tech stack?

- 1 — **Highly Complex:** Multiple legacy systems, poor API support, significant re-engineering needed.
- 2 — **Complex:** Integration is possible but requires significant development effort and risk management.
- 3 — **Moderate:** Integration requires custom work but existing APIs or systems make it achievable.
- 4 — **Manageable:** Clear integration pathway. Existing tools support it. Moderate dev effort.
- 5 — **Simple:** Plug-and-play or low-code integration. Can be live quickly with minimal risk.

Composite Scoring & AI Priority Index

Once each process is scored on all four axes, calculate the AI Priority Index (API) using the formula below. This score determines which processes should be prioritised in the 90-Day Roadmap.

Formula: AI Priority Index = (Process Maturity × 1.0) + (Data Readiness × 1.5) + (Business Impact × 2.0) + (Integration Complexity × 1.5)

Maximum possible score: 30 points | Weights reflect the relative importance of each dimension. Business Impact is weighted highest — a high-impact, complex process outranks a low-impact, simple one.

AI Readiness Tiers — Company-Level Score

The company-level AI Readiness Score is calculated by averaging the AI Priority Index scores across all assessed processes, then normalising to a 100-point scale. This produces the client-facing score delivered on Day 5.

Score Range	Tier	Description	Typical Profile
0 – 33	Starter	Early stage. AI adoption requires foundational work before deployment.	Manual-heavy operations, limited data infrastructure, first-time AI consideration.
34 – 66	Builder	Ready to implement targeted AI solutions in specific, well-defined processes.	Some digitalisation in place, clear pain points, leadership buy-in exists.

67 – 100	Scaler	Strong foundation. Ready for advanced AI deployment and enterprise-grade automation.	Modern tech stack, data-rich, experienced with automation, aiming for competitive advantage.
----------	--------	--	--

How to Reference the Matrix in Client Sessions

- Day 2 intro: "We use the zylo ai readiness Matrix (The ZARM) (The ZARM) (The ZARM) (The ZARM) — a 4-axis scoring model — to assess every process we identify. The four dimensions are Process Maturity, Data Readiness, Business Impact, and Integration Complexity. Each scores 1 to 5. We will walk through each one together."
 - Day 3 presentation: Show the completed scoring table and say: "Based on the Matrix, here are the 12 processes we assessed and how they score."
 - Day 5 presentation cover: Display the client's overall AI Readiness Score prominently. "Your company scores 58/100 on the zylo ai readiness Matrix (The ZARM) (The ZARM) (The ZARM) (The ZARM) — placing you in the Builder tier."
 - In all written deliverables: Every reference to scoring should say "Zylo AI Readiness Matrix™ score" — not just "score".
-

ELEMENT 2 — Pre-Sprint Research Brief Template

Purpose: The Sprint Lead completes this brief before Day 1 of every sprint. It takes 2–3 hours. It is presented at the start of the Day 1 session to demonstrate that Zylo has already done its homework. No boutique firm does this. Every top-tier consulting firm does.

Complete this template for [CLIENT COMPANY NAME] | Sprint Start Date: [DATE] | Sprint Lead: [NAME] | Prepared: [DATE, at least 48 hours before Day 1]

Section A — Company Overview

Company Name: [Full legal name and trading name if different]

Website: [URL]

Founded: [Year]

Headquarters: [City, Country]

Other Offices: [List any known locations]

LinkedIn Headcount: [Number — taken from LinkedIn company page]

LinkedIn Growth (12 months): [Is the team growing, shrinking, or stable?]

Business Model: [B2B / B2C / Marketplace / SaaS / Services / Product — describe in 1-2 sentences]

Revenue Estimate: [From LinkedIn, Crunchbase, news, or company statements if public]

Funding Status: [Bootstrapped / Seed / Series A / PE-backed / Public — source if known]

Section B — Products & Services

[Write a 3-5 sentence summary of what the company actually sells, who their customers are, and what problem they solve. Use the website homepage and About page as primary sources.]

Key Products / Services Identified:

Product/Service 1: [Name and 1-line description]

Product/Service 2: [Name and 1-line description]

Product/Service 3: [Name and 1-line description]

Product/Service 4: [Name and 1-line description]

Section C — Leadership & Decision-Makers

Research the people who will be in the room on Day 1.

Name	Title	LinkedIn Profile	Background Notes	Years at Company
------	-------	------------------	------------------	------------------

[CEO Name]	CEO / Founder	[URL]	[Prior companies, expertise, notable achievements]	[X years]
[COO Name]	COO / Ops Lead	[URL]	[Background, focus areas]	[X years]
[Dept Head]	[Title]	[URL]	[Background]	[X years]

Section D — Technology Stack

Use BuiltWith (builtwith.com), LinkedIn job posts, and the company's website to identify tools in use.

Category	Tool / System Identified	How Identified	Notes
CRM	[e.g. Salesforce / HubSpot / None visible]	[BuiltWith / Job post / Website]	[Version or depth of use if known]
ERP / Ops	[e.g. SAP / NetSuite / Spreadsheets]	[Source]	[Notes]
Communication	[e.g. Slack / Teams / Email-only]	[Source]	[Notes]
Marketing Automation	[e.g. HubSpot / Mailchimp / None]	[Source]	[Notes]
Data / Analytics	[e.g. Tableau / PowerBI / None]	[Source]	[Notes]
Cloud Infrastructure	[e.g. AWS / Azure / GCP / None]	[Source]	[Notes]
Other Notable Tools	[Any other tools visible]	[Source]	[Notes]

Section E — Industry Context & AI Trends

Identify 2-3 AI trends directly relevant to this company's industry. These provide context for Day 1 and strengthen the Sprint Lead's credibility.

Industry: [e.g. B2B SaaS / Logistics / Healthcare / Financial Services / eCommerce]

AI Trend 1: [Name the trend and explain it in 2 sentences. Source the example with a real company.]

AI Trend 2: [Name the trend and explain it in 2 sentences. Source the example with a real company.]

AI Trend 3: [Name the trend and explain it in 2 sentences. Source the example with a real company.]

Relevant Industry Benchmark:

[E.g. "In B2B SaaS, companies of similar size are using AI primarily for customer support automation and churn prediction. The leaders in this space (Intercom, Drift) report 40% reduction in support ticket volume after AI deployment. This will be relevant context for Day 1 if the client has a customer-facing support function."]

Section F — Visible Pain Points & Hypotheses

Based on your research, form 3–5 hypothesis questions to explore on Day 1. These are educated guesses based on what you know about the company. Bringing these to Day 1 signals intelligence and preparation.

Hypothesis 1: [E.g. "Based on their LinkedIn job posts for manual data entry roles, they likely have a reporting or reconciliation bottleneck."]

Hypothesis 2: [E.g. "Their CRM is HubSpot but their ops team is larger than typical — suggests manual handoffs between sales and operations."]

Hypothesis 3: [E.g. "No data analytics tools visible — executive decisions are likely being made on gut feel or delayed reporting."]

Hypothesis 4: [Your observation]

Hypothesis 5: [Your observation]

How to use this in Day 1: On Day 1, open with: "Before we begin, our team spent time researching your business. We have a few early observations we'd like to validate with you — and they may shape where we focus today." Then present 2-3 of your hypotheses as questions. This creates an immediate impression of depth that sets the tone for the entire sprint.

ELEMENT 3 — Visual AI Opportunity Map — Design & Production Guide

What this is: The AI Opportunity Map is currently delivered as an Excel table. This section defines exactly how to produce the one-page visual version that the BA creates after Day 3 and includes in the Day 5 presentation. This visual becomes the most-remembered and most-shared asset from the sprint.

The Quadrant Structure

The visual is a 2x2 priority quadrant. The X-axis represents Business Impact (Low → High). The Y-axis represents Implementation Speed (Slow → Fast). Each AI opportunity is plotted as a labelled circle. Circle size represents the ROI estimate (larger = higher financial return).



Production Instructions for the Business Analyst

1. After Day 3, open the AI Opportunity Map Excel sheet. You have 8–12 scored opportunities.
2. For each opportunity, determine its quadrant position: X-axis = Business Impact score (Axis 3 from the Matrix). Y-axis = Integration Complexity score inverted (Axis 4 — high complexity = slow, so invert the score to represent speed).
3. Open PowerPoint. Use a blank slide (widescreen 16:9).
4. Draw the quadrant using lines and a text box background. Label the four quadrants as shown above.
5. Plot each AI opportunity as a circle (Insert → Shape → Oval). Size the circle proportionally to the ROI estimate (small = under \$20K/year, medium = \$20-50K, large = \$50K+).
6. Label each circle with the opportunity name (e.g. "Invoice Automation", "Support Bot", "Sales Forecasting").
7. Colour-code by AI type: Blue = AI Agent, Green = Workflow Automation, Orange = Predictive Analytics, Purple = Document Intelligence.
8. Add Zylo logo, client name, and date in the footer. Title: "[Client Name] — AI Opportunity Map | Zylo AI Readiness Sprint | [Month Year]".
9. Export as PDF. This is Deliverable 3 — the standalone visual version alongside the Excel.

10. Insert this visual as the first content slide of the Day 5 presentation (after the cover and score slide).

The Naming Convention for Opportunities

Every opportunity on the map must have a plain-language name — not a technical one. The client's CEO must be able to understand what each bubble represents at a glance.

Technical Name to Avoid	Plain-Language Name to Use
NLP-based ticket classification system	Smart Support Inbox
Predictive churn ML model	At-Risk Customer Alert
Invoice OCR and ERP integration	Automated Invoice Processing
LLM-powered outbound generation	AI Sales Outreach Assistant
Automated KPI dashboard pipeline	Live Business Dashboard
Demand forecasting regression model	Inventory Demand Predictor

ELEMENT 4 — AI Readiness Score Card — System & Templates

What this is: Every client receives a branded AI Readiness Score as part of their Day 5 deliverables. This section defines how the score is calculated, how it is presented, and the exact template text for the LinkedIn-shareable certificate.

Score Calculation Formula

The AI Readiness Score is a normalised composite of the zylo ai readiness Matrix (The ZARM) results across all assessed processes. Follow these steps:

11. Score every identified process across all 4 axes (1–5 each).
12. Calculate the AI Priority Index for each process using the weighted formula: $(\text{Process Maturity} \times 1.0) + (\text{Data Readiness} \times 1.5) + (\text{Business Impact} \times 2.0) + (\text{Integration Complexity} \times 1.5)$.
13. Calculate the company's Average AI Priority Index: $\text{sum all process scores} \div \text{number of processes assessed}$.
14. Normalise to 100: $(\text{Average API} \div 30) \times 100 = \text{AI Readiness Score}$.
15. Round to nearest whole number. Assign the tier: 0–33 = Starter, 34–66 = Builder, 67–100 = Scaler.

Worked Example: Example: Company assessed 10 processes. Total API = 185. Average API = 18.5. Normalised Score = $(18.5 \div 30) \times 100 = 62$. Tier: Builder.

Score Card — Day 5 Presentation Slide Template

The first content slide of every Day 5 presentation must display the score prominently. Here is the exact copy template:

[CLIENT COMPANY NAME]

AI Readiness Score

[SCORE] / 100

Tier: [STARTER / BUILDER / SCALER]

"[Company Name] has been assessed across [X] business processes using the zylo ai readiness Matrix (The ZARM). With a score of [X]/100, you are positioned in the [Tier] tier — meaning [tier description]. This sprint identified [X] high-priority AI opportunities that, if implemented, would move your score to an estimated [projected score] within 12 months."

Assessed by Zylo Solutions | [Month Year]

Score Card — PDF Certificate Template (Delivered Post-Sprint)

After the sprint, the BA generates a one-page branded PDF certificate for the client to use internally and share on LinkedIn. Here is the exact content:

ZYLO AI READINESS CERTIFICATE

This certifies that

[COMPANY NAME]

has completed the Zylo AI Readiness Sprint™
and achieved an AI Readiness Score of

[SCORE] / 100 — [TIER] Tier

[TIER DESCRIPTION — 1 sentence]

Assessed across [X] business processes using the zylo ai readiness Matrix (The ZARM)

Issued by Zylo Solutions | [Date]
[Sprint Lead Name], AI Strategy Consultant

zylosolutions.com

LinkedIn Post Template for Clients (Send to Client to Encourage Sharing)

Subject: A quick ask — would you be open to sharing your AI Readiness Score?

Hi [Name],

We have attached your AI Readiness Certificate as a PDF. Some of our clients find it useful to share their score on LinkedIn — it starts great conversations and signals to their team and partners that they are taking AI seriously.

Here is a suggested post caption if you would like to use it:

"We just completed a 5-day AI Readiness Sprint with Zylo Solutions and received our company's AI Readiness Score: [X]/100 — placing us in the [Tier] tier. Came away with a full AI Opportunity Map and a 90-day implementation roadmap. If you're curious about where your business stands on AI readiness, I'd recommend exploring it. [Tag Zylo Solutions]"

No pressure at all — just sharing in case it's useful.

[Rep Name]

ELEMENT 5 — Sprint Retrospective — 30-Minute Internal Agenda & Template

Rules: This retrospective is mandatory within 48 hours of every Day 5 delivery. Attendees: Sprint Lead, Business Analyst, Technical Architect, Rep. It is not a debrief — it is a structured improvement session. Findings feed into the next version of the playbook.

Agenda (30 Minutes)

0–5 min: Context recap — Rep shares: client profile, industry, company size, score delivered, and initial conversion signal (Did they say yes / not yet / no on Day 5?)

5–10 min: What worked well? — Each team member names ONE thing that went better than expected. Sprint Lead records these.

10–18 min: What did not work? — Open discussion. Specific only — no vague feedback. Each issue must be stated as: "On [Day X], [what happened] — the client's reaction was [X] — we should have done [Y] instead."

18–24 min: What did we miss? — What did the client say or reveal that we did not capture in the deliverables? Any signals we ignored that we should have built on?

24–28 min: One change for next sprint — The team agrees on ONE specific change to make to the framework, workbook, presentation, or process before the next sprint.

28–30 min: Conversion status & next step — Rep confirms what follow-up action is scheduled and by when.

Retrospective Record Sheet — Fill In After Every Sprint

Sprint Number: [#]

Client Company: [Name]

Client Industry: [Industry]

AI Readiness Score Delivered: [X/100 | Tier]

Day 5 Conversion Signal: [Yes — Proposal Requested / Not Yet — Following Up / No — Reasons Below]

Sprint Lead: [Name]

Date of Retrospective: [Date]

What Worked Well

1. [Record specific thing that worked — who did it, when, what the impact was]
2. [Record specific thing that worked — who did it, when, what the impact was]
3. [Record specific thing that worked — who did it, when, what the impact was]

What Did Not Work

- Issue 1:** [Day X — what happened — client reaction — what to do differently]
- Issue 2:** [Day X — what happened — client reaction — what to do differently]
- Issue 3:** [Day X — what happened — client reaction — what to do differently]

What We Missed

- Client insight not in deliverables:** [What the client said that we did not capture or build on]
- Signal we ignored:** [Any buying or objection signal we noticed but did not address]

One Change for the Next Sprint

[Write the specific, actionable change agreed upon. E.g. "From Sprint 4 onwards, the BA will prepare a draft Process Inventory list from the company's job posts before Day 1 — to be validated on Day 2 rather than built from scratch."]

Cumulative Learning Log (Update After Every Sprint)

Sprint #	Key Learning	Change Made	Impact Observed
[1]	[What was learned]	[What was changed]	[What improved]
[2]			
[3]			

ELEMENT 6 — Sprint Debrief Loom Script — Post-Day 5

Personal Video

What this is: This 5–7 minute personal video is recorded by the Sprint Lead on the evening of Day 5 and sent to the client the following morning. It is NOT in the deliverables package — it is unexpected. It is the single most-memorable thing Zylo does. Record it on Loom (free tier works fine). Do not over-produce it. Sincerity beats polish here.

Recording Rules: Do: Record in one take if possible. Sit in front of a clean background. Speak directly to camera. Reference the client by name. Be specific — reference actual findings from their sprint. Aim for 5-7 minutes.

Do NOT: Write a full script and read it. Over-edit. Use slides. Make it a sales pitch. Reference the proposal.

Script Framework — Personalise All [Bracketed Sections]

Opening (30 seconds)

"Hi [Name] — [Sprint Lead Name] here. It's the evening after Day 5 and I just wanted to record a quick personal note for you — separate from the formal deliverables we sent over. This is just me, off the back of a week with your business, sharing what genuinely stood out."

Observation 1 — The Most Important Finding (90 seconds)

"The thing that struck me most about [Company Name] — and I've done a lot of these sprints — is [specific observation, e.g. "the gap between how well your sales team performs and how slow your back-office operation is moving. That bottleneck is costing you significantly, and it's entirely solvable."]

The reason this matters more than the other opportunities on the map is [explanation]. I've seen this pattern in [similar industry] companies before, and the ones who address it first tend to pull ahead of their competitors meaningfully within 12 months."

Observation 2 — The Hidden Opportunity (90 seconds)

"The second thing I want to call out — and this didn't make it to the top of the roadmap scoring because the data isn't quite there yet — is [specific observation, e.g. "your customer data. You are sitting on [X years] of transaction history that is currently just sitting in [system]. Once we build the infrastructure to use it, the predictive capability it unlocks is significant."]

I mention this not because it's in Phase 1, but because it's worth thinking about as you build your data infrastructure in the next 6 months. It will pay off later."

Observation 3 — What I Would Prioritise If It Were My Business (60 seconds)

"If I were in [Name]'s position — honestly — I would start with [specific Phase 1 opportunity]. Not because it has the highest score on the matrix, but because it builds a foundation that every other AI system you add later will depend on. Get that right first, and everything else gets easier."

Closing (45 seconds)

"I genuinely enjoyed this sprint. Your team was [specific compliment — e.g. "unusually open about the real challenges, which made the work much more useful than it could have been"]. The deliverables are in your inbox. [Rep Name] will follow up with you — no pressure at all. But if you have any questions between now and then, feel free to reply directly to me.

Good luck with it. I hope this week was useful."

Sending Instructions

16. Record on Loom (loom.com — free account is fine).
 17. Title the video: "[Client Company] — AI Readiness Sprint | Personal Note from [Sprint Lead Name]"
 18. Copy the Loom share link.
 19. Send via email the morning after Day 5 (not the same night — let it arrive fresh).
 20. Subject line: "A personal note from [Sprint Lead Name] — after this week with your team"
 21. Email body: "Hi [Name], I recorded a short personal note last night — separate from the deliverables. Wanted to share a few things I didn't get to say fully in the session. [LOOM LINK]. No ask, no pitch — just a few honest observations. [Sprint Lead Name]"
-

ELEMENT 7 — Reference Client Programme

What this is: After Sprint 3, every converted client is formally invited into the Zylo Reference Programme. A reference call from a peer company closes faster than any proposal. This programme should be built into the post-project relationship from day one — not added as an afterthought.

The Ask — When and How to Ask

Ask for reference status at the end of the first successful project milestone — not at the end of the sprint (too early) and not after 6 months (too late to feel meaningful).

Reference Ask Script: "[Name], we've hit a strong milestone together — [specific result, e.g. "the invoice automation is live and saving your team 14 hours a week"]. We have a small programme where satisfied clients occasionally take a short 15-minute call with a company we're about to work with — someone who has similar challenges and wants to hear a peer's honest view before committing. We never pressure them, we always ask your permission first, and it never takes more than 15 minutes.

Would you be comfortable being part of that? We'd love to have your voice in the room for the right conversation."

What Zylo Offers Reference Clients in Return

- Monthly 30-minute AI Strategy Check-In call with the Sprint Lead — no agenda, just an ongoing conversation about what's working and what's new in AI relevant to their business.
- Access to Zylo's internal resource library — new frameworks, tools, and research shared before it goes public.
- Early access to any new Zylo services or tools at preferential rates.
- Introduction to other Zylo clients in compatible (non-competing) industries when relevant.

Reference Client Registry — Track All Reference Clients

Client Name	Company	Industry	Project Delivered	Key Outcome Achieved	Contact Preference	Times Used	Last Used
[Name]	[Company]	[Industry]	[Project type]	[Specific result]	[Call / Email]	[#]	[Date]
[Name]	[Company]	[Industry]	[Project type]	[Specific result]	[Call / Email]	[#]	[Date]

Reference Call Briefing — What to Send the Reference Client Before the Call

Subject: Reference Call Brief — [Prospect Company Name]

Hi [Reference Client Name],

Thank you for agreeing to take this call. Here is a brief on the company you will be speaking with:

Company: [Prospect Company Name]

Industry: [Industry]

Size: [Headcount / Revenue range]

Challenge they are exploring: [1-2 sentences — what they want to solve with AI]

Why we matched you: [Specific reason — similar industry, similar size, similar challenge]

The call will be 15 minutes. They may ask:

- What the sprint process was like
- Whether the deliverables were useful
- What the working relationship with Zylo is like
- Whether the outcome matched the promise

You do not need to prepare anything — just speak honestly. That is why it works.

Thank you again. [Rep Name]

ELEMENT 8 — Technical FAQ — 20 Questions Every Sprint Lead Must Know

How to use this: Every Sprint Lead must be able to answer these 20 questions confidently in a client session. If a question is beyond scope, the protocol is: "Great question — I want to give you a precise answer rather than a rough one. Let me bring in our Technical Architect to confirm the specifics and get back to you within 24 hours." Never guess. Never bluff.

Category 1 — AI Agents

Q1: What exactly is an AI agent and how is it different from basic automation?

Answer: Basic automation follows fixed rules — if X happens, do Y. An AI agent can reason, make decisions, and handle situations it has not seen before. It uses a language model or ML model to interpret context and choose the appropriate action. Example: a basic automation routes emails by keyword. An AI agent reads the email, understands the intent, drafts a personalised response, and escalates if uncertain.

Q2: Can an AI agent work inside our existing tools (CRM, ERP, Slack)?

Answer: Yes. AI agents connect to existing tools via APIs. Most modern SaaS tools (Salesforce, HubSpot, Slack, Notion, SAP, etc.) have APIs. The agent is built to read from and write to these systems. The key question is not whether it can connect — it usually can — but how clean and accessible the data in those tools is.

Q3: How long does it take to build an AI agent?

Answer: A simple AI agent (e.g. a customer support triage bot) can be built in 2–4 weeks. A more complex agent (e.g. one that reads contracts, extracts data, and updates a CRM) typically takes 6–10 weeks. The timeline depends on data readiness, integration complexity, and how many edge cases need to be handled.

Q4: What happens when the AI agent makes a mistake?

Answer: AI agents make mistakes — this is expected and managed by design. Best practice is to build in a "confidence threshold" — if the agent's confidence in its output is below a set level, it flags the case for human review rather than acting autonomously. Over time, flagged cases are used to improve the model. Most clients run agents in a "shadow mode" first — it runs in parallel, a human checks its outputs before they are actioned.

Q5: Do we own the AI agent or does Zylo?

Answer: You own everything Zylo builds for you. All code, models, and systems are delivered to you as part of the project. Zylo may retain the right to use anonymised, aggregated learnings to improve methodology — this is covered in the service agreement — but the specific system built for your business belongs to you.

Category 2 — Workflow Automation

Q6: What is the difference between AI automation and RPA (Robotic Process Automation)?

Answer: RPA mimics human actions on a screen — clicking, copying, pasting — without understanding what it is doing. It breaks when the interface changes. AI automation understands the content and intent of data, not just its position on a screen. AI automation is more resilient, more flexible, and more valuable for complex workflows. Zylo builds AI automation — not RPA.

Q7: Which processes are best suited for automation?

Answer: Processes with high volume, clear rules (or learnable patterns), structured data input, and repetitive steps. Prime examples: invoice processing, report generation, data entry between systems, appointment scheduling, contract review for standard clauses, and customer query routing.

Q8: How do we measure the ROI of an automation?

Answer: The standard formula: $(\text{Hours saved per week} \times \text{hourly cost} \times 52 \text{ weeks}) + (\text{Error reduction value}) + (\text{Speed improvement value}) - (\text{Build cost} + \text{annual maintenance})$. We model this during Day 4 using your actual team cost data. Most SMB automations achieve full ROI within 3–9 months.

Q9: What tools does Zylo use to build automations?

Answer: Depends on the complexity and client tech stack. Common tools include: n8n and Make (Integromat) for workflow orchestration, OpenAI / Anthropic APIs for language understanding, Python for custom logic, and existing client APIs for integration. We choose the right tool for the specific use case — not a one-size-fits-all approach.

Q10: What ongoing maintenance does an automation require?

Answer: Most automations require light monthly maintenance — primarily monitoring for errors and updating when connected systems change their APIs. Zylo provides a maintenance recommendation as part of every project delivery. Typically 2–4 hours per month for a standard automation.

Category 3 — Machine Learning & Predictive Analytics

Q11: Do we need a data science team to use machine learning?

Answer: No. Zylo builds and deploys the ML model for you. What you need is: sufficient historical data (typically 1–2 years of relevant records), a clear question you want the model to answer (e.g. "which customers are most likely to churn this month?"), and a team member who can interpret and act on the model's output. The data science work is done by Zylo.

Q12: How much data do we need to train a machine learning model?

Answer: It varies by use case. As a general guide: classification models (e.g. spam detection, churn prediction) need a minimum of 1,000–5,000 labelled examples. Forecasting models need 1–2 years of historical data with consistent frequency. Document intelligence models can work with as few as 200–500 labelled samples. We assess your data readiness during Day 2 of the sprint and flag any data gaps clearly.

Q13: How accurate will the predictions be?

Answer: Accuracy depends on data quality and complexity of the prediction task. Typical real-world accuracy for SMB use cases: churn prediction 75–85%, demand forecasting 80–90% (within a defined range), document classification 90–95%. We always show the model's accuracy and uncertainty range as part of delivery — and set client expectations realistically during Day 4.

Q14: Can machine learning models be biased? How do we address that?

Answer: Yes. ML models learn patterns from historical data — if that data reflects biased decisions, the model will replicate them. Zylo conducts a bias audit for any model that makes decisions about people (hiring, credit, customer treatment). For operational models (demand forecasting, process routing), bias is less of a concern but we still review training data for anomalies.

Q15: What happens to the model over time — does it get worse?

Answer: Yes — this is called model drift. As business conditions change, a model trained on old data becomes less accurate. Best practice is to retrain the model quarterly or when accuracy drops below a defined threshold. Zylo builds a monitoring dashboard into every ML delivery that alerts when accuracy degrades.

Category 4 — Data, Security & Compliance

Q16: Is our data safe when we share it with Zylo?

Answer: Yes. All client data is covered by the NDA signed before the sprint begins. Zylo does not store, share, or use client data for any purpose outside the engagement. If we use cloud AI APIs (OpenAI, Anthropic, etc.), we use their enterprise tiers with data privacy agreements in place. We can provide a full data handling policy on request.

Q17: Can AI be used if we have GDPR / data compliance requirements?

Answer: Yes, with the right architecture. AI systems can be designed to be GDPR-compliant: using anonymised or pseudonymised data, keeping data within approved geographic regions (EU servers), implementing right-to-erasure workflows, and maintaining data processing records. Zylo flags compliance requirements during the Day 2 assessment and factors them into every recommendation.

Q18: Do we need to move our data to the cloud to use AI?

Answer: Not necessarily. Many AI solutions can be deployed on-premise or in a private cloud environment. However, cloud-based deployment is significantly faster and more cost-effective for most SMBs. We discuss infrastructure requirements during Day 2 and recommend the deployment model that balances capability, cost, and compliance for your specific situation.

Q19: What if our data quality is poor? Can we still use AI?

Answer: Poor data quality is one of the most common barriers we find — and it is addressable. The first step is a data quality audit (part of Day 2). For processes where data quality is low, we recommend a "data pipeline" as the first build — a system that cleans, structures, and enriches your data going forward. This is usually a 2–4 week project that unlocks all subsequent AI use cases.

Q20: How long until we see results?

Answer: Quick Win opportunities (low complexity, high data readiness) can be live within 4–8 weeks and deliver measurable results within 30 days of going live. Foundation builds typically take 8–14 weeks. Strategic investments (complex integrations, ML models) take 3–6 months. The 90-Day Roadmap is specifically designed to deliver a live, measurable result within the first 30 days — so the client sees real value before the relationship is even fully established.

ELEMENT 9 — Industry Benchmark Template — 90-Day Roadmap Insert

What this is: This one-page section is added to the 90-Day Roadmap document as Page 2, before the roadmap itself. It gives the client market context — how they compare to peers — and contextualises the urgency. The Sprint Lead researches and fills this in during Day 3/4. It takes approximately 1 hour.

Benchmark Section — Template Copy (Insert Client Data)

INDUSTRY CONTEXT — WHERE [COMPANY NAME] STANDS

AI Adoption in [Industry] — [Month Year]

Companies in [Industry] at a similar stage to [Company Name] ([headcount range], [revenue range]) are currently at [adoption level] of AI maturity on average. Here is how the market breaks down:

- The Leaders ([top 15% of industry]): Already operating with [specific AI use — e.g. "fully automated customer support triage, predictive inventory management, and AI-assisted sales outreach"]. Representative companies: [2-3 named examples if publicly known].
- The Mainstream ([middle 60%]): Actively experimenting with AI in [1-2 specific areas]. Most have deployed at least one automation but have not built a systematic AI strategy.
- The Laggards ([bottom 25%]): Still operating primarily on manual processes. At risk of being outcompeted as AI-enabled peers reduce costs and accelerate speed.

Where [Company Name] sits today:

Based on our assessment, [Company Name] is currently in the [Leader / Mainstream / Laggard] segment, scoring [X]/100 on the zylo ai readiness Matrix (The ZARM) . [One sentence on the specific gap or advantage this represents.]

Where this roadmap takes you:

If Phase 1 and Phase 2 are implemented within the 90-day timeline, we project your AI Readiness Score will move from [current score] to approximately [projected score] — moving you from [current tier] to [target tier] and placing you ahead of [X%] of your industry peers.

Research Sources — Where to Find Industry Benchmark Data (1 Hour Research)

Source	What to Find There	URL / Access
McKinsey Global Institute Reports	AI adoption rates by industry, productivity impact studies	mckinsey.com/featured-insights

Gartner Hype Cycle Reports	AI maturity by technology type and sector	gartner.com/en/research (free summaries)
IBM Global AI Adoption Index	Annual survey of AI adoption by industry and region	ibm.com/thought-leadership/institute-business-value
Deloitte AI Institute	Industry-specific AI use case adoption data	deloitte.com/us/en/insights/deloitte-ai-institute
LinkedIn Industry Reports	Hiring trends revealing AI tool adoption in specific industries	linkedin.com/pulse/topics
Google Trends	Search volume for AI terms by industry — signals adoption momentum	trends.google.com

Industry Benchmark Snapshots — Pre-Researched for Common Zylo Target Industries

B2B SaaS

AI adoption is high in sales and marketing (AI SDRs, lead scoring) and customer success (churn prediction). Leaders are deploying AI for product feature development (co-pilots, recommendations). Average companies are using AI for support automation. Laggards are still manually managing customer journeys.

Estimated AI adoption rate in this sector: 68%

Most common use cases: Support ticket automation, sales lead scoring, churn prediction

eCommerce / Retail

Leaders are using AI for personalised product recommendations, dynamic pricing, and demand forecasting. Mainstream companies are deploying chatbots and basic inventory alerts. The gap between leaders and laggards is widening fastest in this sector.

Estimated AI adoption rate in this sector: 72%

Most common use cases: Product recommendations, inventory forecasting, customer segmentation

Professional Services (Consulting, Legal, Accounting)

Adoption is slower due to regulatory caution but accelerating. Leaders are using AI for document review, contract analysis, and proposal generation. Mainstream firms are experimenting with AI research assistants.

Estimated AI adoption rate in this sector: 41%

Most common use cases: Document intelligence, proposal automation, research assistance

Logistics & Supply Chain

AI is predominantly used for route optimisation, demand forecasting, and warehouse automation in large firms. SMBs in this sector are significantly behind — most are still on spreadsheets for planning.

Estimated AI adoption rate in this sector: 55%

Most common use cases: Route optimisation, demand forecasting, shipment tracking automation

Healthcare (Private / Non-clinical)

AI for appointment scheduling, administrative automation, and patient communication is mainstream. Clinical AI is heavily regulated. Non-clinical operations are the safest and fastest area for AI deployment.

Estimated AI adoption rate in this sector: 48%

Most common use cases: Appointment automation, admin workflows, patient communication

Financial Services (SMB)

Compliance-heavy environment but AI for fraud detection, report generation, and client communication is well-established in larger firms. SMBs are adopting AI reporting tools rapidly.

Estimated AI adoption rate in this sector: 59%

Most common use cases: Automated reporting, client communication, document processing

ELEMENT 10 — Zylo AI Leaders Network — Alumni Community Playbook

What this is: The Zylo AI Leaders Network is a private community for companies that have completed the AI Readiness Sprint. It positions Zylo as the hub of an expert community — not just a vendor. It creates ongoing touchpoints, peer referrals, and inbound enquiries with zero ad spend.

Setup — What to Create (Before Sprint 3)

22. Create a private LinkedIn group: Name = "Zylo AI Leaders Network | Members Only". Description: "A private community for leaders who have completed the Zylo AI Readiness Sprint — sharing AI implementation experiences, lessons, and opportunities." Set to private/request-only.
23. Create a WhatsApp Community (parallel to LinkedIn for GCC and mobile-first markets). Name: "Zylo AI Leaders Network". Admin: Sprint Lead or Rep.
24. Create a simple welcome document (1 page): community purpose, rules, what to expect, and how to get help from other members.
25. Designate one Zylo team member as Community Manager — responsible for posting weekly content and monitoring engagement. Time commitment: 30 minutes per week.

Invitation — When and How to Invite Sprint Alumni

Invitation Message: Subject: An invitation — Zylo AI Leaders Network

Hi [Name],

We have built a private community for leaders who have completed the Zylo AI Readiness Sprint — a space to share what is working, what is not, and what is coming next in AI.

The community is deliberately small and selective — only sprint alumni are invited. No vendors, no sales pitches. Just a group of business leaders navigating similar challenges.

We post one genuinely useful AI insight per week and occasionally facilitate introductions between members with complementary businesses.

If you would like to join:

- LinkedIn Group: [LINK]
- WhatsApp Community: [LINK]

No obligation. But if you find value in it, we would love to have you in the room.

[Rep Name] | Zylo Solutions

Weekly Content Calendar — What to Post (30 Minutes Per Week)

Week	Content Type	Example Topic	Format
Week 1 of month	AI Tool Spotlight	"The AI tool we're using in client sprints right now — and why"	Short text post + 1 screenshot
Week 2 of month	Industry Insight	"What AI adoption looks like in [industry] right now — 3 observations"	Text post with data point
Week 3 of month	Case Study / Story	"How one of our clients saved X hours a week with [automation type]" (anonymised or with permission)	Story-format post
Week 4 of month	Member Question / Discussion	"Question for the group: what's the biggest bottleneck slowing your AI implementation right now?"	Discussion prompt

Community Rules — Post in the Group on Day 1

- This is a peer-to-peer knowledge community. No vendor pitches or promotional content.
- Members may introduce themselves and their business once.
- Questions are welcomed and encouraged — someone in the group has likely solved your problem.
- Zylo posts weekly insights but does not dominate the conversation — this is your community.
- Any introductions between members must be requested and consented to by both parties.

How the Community Generates Business for Zylo

The community generates business through three mechanisms — none of which involve direct selling:

26. Peer referrals: Member A mentions Zylo in conversation with a non-member peer. The community gives them a specific thing to reference: "They run a community for alumni — it's worth applying for the sprint just to get access."
27. Inbound from member activity: When members post about their AI implementations, their networks see it. Some of those viewers will enquire with Zylo directly.
28. Upsell conversations: Monthly, Zylo reviews which community members are engaging most. These are warm contacts for deeper conversations about Phase 2 or 3 implementation — no cold outreach needed.

Milestone Goals for the Community

Milestone	Target	Action When Reached
First 10 members	After Sprint 3–4	Send personal welcome message to each. Ask first discussion question to activate the group.

First 25 members	After Sprint 6–8	Host first live virtual event: 45-min "AI Implementation Q&A" with Sprint Lead. Record and share.
First 50 members	After Sprint 12–15	Launch monthly "AI Leader Spotlight" — interview a member about their AI journey. Share in community and on Zylo LinkedIn.
First 100 members	After Sprint 25–30	Consider an annual in-person event in primary market city. Community is now a competitive moat.